

```
In [51]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns

import warnings
warnings.filterwarnings('ignore')

## Sklearn Library
from scipy.stats import uniform
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, roc_auc_score
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import confusion_matrix
from sklearn.metrics import matthews_corrcoef
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
import sklearn.metrics as sm
```

```
In [2]: df=pd.read_csv(r"C:\Users\DD\OneDrive\Desktop\ML PROJECTS\credit fraud\creditcard.csv")
```



In [3]: df

Out[3]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	...	-0.018307	0.277
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	...	-0.225775	-0.638
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	...	0.247998	0.771
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	...	-0.108300	0.005
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	...	-0.009431	0.798
...	...	...	...	...	...	...	...	...	...	...	...	...	...
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	-2.606837	-4.918215	7.305334	1.914428	...	0.213454	0.111
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	1.058415	0.024330	0.294869	0.584800	...	0.214205	0.924
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	3.031260	-0.296827	0.708417	0.432454	...	0.232045	0.578
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	0.623708	-0.686180	0.679145	0.392087	...	0.265245	0.800
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	-0.649617	1.577006	-0.414650	0.486180	...	0.261057	0.643

284807 rows × 31 columns



In [4]: df.info



```

Out[4]: <bound method DataFrame.info of
0      0.0  -1.359807  -0.072781  2.536347  1.378155  -0.338321
1      0.0   1.191857   0.266151  0.166480  0.448154  0.060018
2      1.0  -1.358354  -1.340163  1.773209  0.379780  -0.503198
3      1.0  -0.966272  -0.185226  1.792993  -0.863291  -0.010309
4      2.0  -1.158233   0.877737  1.548718  0.403034  -0.407193
...
284802 172786.0 -11.881118 10.071785 -9.834783 -2.066656 -5.364473
284803 172787.0  -0.732789  -0.055080  2.035030  -0.738589  0.868229
284804 172788.0  1.919565  -0.301254  -3.249640  -0.557828  2.630515
284805 172788.0  -0.240440   0.530483  0.702510  0.689799  -0.377961
284806 172792.0  -0.533413  -0.189733  0.703337  -0.506271  -0.012546

      V6      V7      V8      V9  ...      V21      V22  \
0      0.462388  0.239599  0.098698  0.363787  ...  -0.018307  0.277838
1     -0.082361 -0.078803  0.085102 -0.255425  ...  -0.225775 -0.638672
2      1.800499  0.791461  0.247676 -1.514654  ...   0.247998  0.771679
3      1.247203  0.237609  0.377436 -1.387024  ...  -0.108300  0.005274
4      0.095921  0.592941 -0.270533  0.817739  ...  -0.009431  0.798278
...
284802 -2.606837 -4.918215  7.305334  1.914428  ...   0.213454  0.111864
284803  1.058415  0.024330  0.294869  0.584800  ...   0.214205  0.924384
284804  3.031260 -0.296827  0.708417  0.432454  ...   0.232045  0.578229
284805  0.623708 -0.686180  0.679145  0.392087  ...   0.265245  0.800049
284806 -0.649617  1.577006 -0.414650  0.486180  ...   0.261057  0.643078

      V23      V24      V25      V26      V27      V28  Amount  \
0     -0.110474  0.066928  0.128539 -0.189115  0.133558 -0.021053  149.62
1      0.101288 -0.339846  0.167170  0.125895 -0.008983  0.014724   2.69
2      0.909412 -0.689281 -0.327642 -0.139097 -0.055353 -0.059752  378.66
3     -0.190321 -1.175575  0.647376 -0.221929  0.062723  0.061458  123.50
4     -0.137458  0.141267 -0.206010  0.502292  0.219422  0.215153   69.99
...
284802  1.014480 -0.509348  1.436807  0.250034  0.943651  0.823731   0.77
284803  0.012463 -1.016226 -0.606624 -0.395255  0.068472 -0.053527  24.79
284804 -0.037501  0.640134  0.265745 -0.087371  0.004455 -0.026561  67.88
284805 -0.163298  0.123205 -0.569159  0.546668  0.108821  0.104533  10.00
284806  0.376777  0.008797 -0.473649 -0.818267 -0.002415  0.013649  217.00

      Class
0         0
1         0
2         0

```



```
3      0
4      0
...    ...
284802  0
284803  0
284804  0
284805  0
284806  0
```

```
[284807 rows x 31 columns]>
```



In [6]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Time        284807 non-null float64
1   V1          284807 non-null float64
2   V2          284807 non-null float64
3   V3          284807 non-null float64
4   V4          284807 non-null float64
5   V5          284807 non-null float64
6   V6          284807 non-null float64
7   V7          284807 non-null float64
8   V8          284807 non-null float64
9   V9          284807 non-null float64
10  V10         284807 non-null float64
11  V11         284807 non-null float64
12  V12         284807 non-null float64
13  V13         284807 non-null float64
14  V14         284807 non-null float64
15  V15         284807 non-null float64
16  V16         284807 non-null float64
17  V17         284807 non-null float64
18  V18         284807 non-null float64
19  V19         284807 non-null float64
20  V20         284807 non-null float64
21  V21         284807 non-null float64
22  V22         284807 non-null float64
23  V23         284807 non-null float64
24  V24         284807 non-null float64
25  V25         284807 non-null float64
26  V26         284807 non-null float64
27  V27         284807 non-null float64
28  V28         284807 non-null float64
29  Amount      284807 non-null float64
30  Class       284807 non-null int64
dtypes: float64(30), int64(1)
memory usage: 67.4 MB
```



```
In [7]: df.isnull()
```

Out[7]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22	V23	V24	V25	V26	V27	V28	Am
0	False	False	False	False	False	False	False	False	False	False	...	False	False	False	False	False	False	False	False	f
1	False	False	False	False	False	False	False	False	False	False	...	False	False	False	False	False	False	False	False	f
2	False	False	False	False	False	False	False	False	False	False	...	False	False	False	False	False	False	False	False	f
3	False	False	False	False	False	False	False	False	False	False	...	False	False	False	False	False	False	False	False	f
4	False	False	False	False	False	False	False	False	False	False	...	False	False	False	False	False	False	False	False	f
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	
284802	False	False	False	False	False	False	False	False	False	False	...	False	False	False	False	False	False	False	False	f
284803	False	False	False	False	False	False	False	False	False	False	...	False	False	False	False	False	False	False	False	f
284804	False	False	False	False	False	False	False	False	False	False	...	False	False	False	False	False	False	False	False	f
284805	False	False	False	False	False	False	False	False	False	False	...	False	False	False	False	False	False	False	False	f
284806	False	False	False	False	False	False	False	False	False	False	...	False	False	False	False	False	False	False	False	f

284807 rows × 31 columns



```
In [8]: df.isnull().sum()
```

```
Out[8]: Time      0
V1      0
V2      0
V3      0
V4      0
V5      0
V6      0
V7      0
V8      0
V9      0
V10     0
V11     0
V12     0
V13     0
V14     0
V15     0
V16     0
V17     0
V18     0
V19     0
V20     0
V21     0
V22     0
V23     0
V24     0
V25     0
V26     0
V27     0
V28     0
Amount   0
Class    0
dtype: int64
```

```
In [11]: df['Class'].value_counts()
```

```
Out[11]: Class
0      284315
1        492
Name: count, dtype: int64
```





```
In [12]: # separating the data for analysis
legit = df[df.Class == 0]
fraud = df[df.Class == 1]
```

```
In [13]: legit.head()
```

Out[13]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22	
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	...	-0.018307	0.277838	-0.110
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	...	-0.225775	-0.638672	0.101
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	...	0.247998	0.771679	0.909
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	...	-0.108300	0.005274	-0.190
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	...	-0.009431	0.798278	-0.137

5 rows × 31 columns



```
In [14]: fraud.head()
```

Out[14]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22	
541	406.0	-2.312227	1.951992	-1.609851	3.997906	-0.522188	-1.426545	-2.537387	1.391657	-2.770089	...	0.517232	-0.035049	-0.110
623	472.0	-3.043541	-3.157307	1.088463	2.288644	1.359805	-1.064823	0.325574	-0.067794	-0.270953	...	0.661696	0.435477	0.101
4920	4462.0	-2.303350	1.759247	-0.359745	2.330243	-0.821628	-0.075788	0.562320	-0.399147	-0.238253	...	-0.294166	-0.932391	0.909
6108	6986.0	-4.397974	1.358367	-2.592844	2.679787	-1.128131	-1.706536	-3.496197	-0.248778	-0.247768	...	0.573574	0.176968	-0.190
6329	7519.0	1.234235	3.019740	-4.304597	4.732795	3.624201	-1.357746	1.713445	-0.496358	-1.282858	...	-0.379068	-0.704181	-0.137

5 rows × 31 columns



```
In [15]: # statistical measures of the data
legit.Amount.describe()
```

```
Out[15]: count      284315.000000
         mean         88.291022
         std        250.105092
         min           0.000000
         25%          5.650000
         50%         22.000000
         75%         77.050000
         max       25691.160000
         Name: Amount, dtype: float64
```

```
In [16]: fraud.Amount.describe()
```

```
Out[16]: count         492.000000
         mean        122.211321
         std        256.683288
         min           0.000000
         25%          1.000000
         50%          9.250000
         75%       105.890000
         max       2125.870000
         Name: Amount, dtype: float64
```

```
In [17]: # compare the values for both transactions
df.groupby('Class').mean()
```

```
Out[17]:
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V20	
Class													
0	94838.202258	0.008258	-0.006271	0.012171	-0.007860	0.005453	0.002419	0.009637	-0.000987	0.004467	...	-0.000644	-0.000644
1	80746.806911	-4.771948	3.623778	-7.033281	4.542029	-3.151225	-1.397737	-5.568731	0.570636	-2.581123	...	0.372319	0.700000

2 rows × 30 columns



# sample

```
In [19]: legit_sample = legit.sample(n=800)
```

```
In [20]: new_dataset = pd.concat([legit_sample, fraud], axis=0)
```

```
In [21]: new_dataset
```

```
Out[21]:
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V
<b>256222</b>	157591.0	-2.456527	2.457377	-0.278652	-1.334673	-0.388679	-0.391650	-0.056992	0.619518	1.626584	...	-0.380373	-0.6583
<b>221908</b>	142745.0	2.192551	-0.697003	-2.172062	-0.869045	-0.185080	-1.774343	0.265166	-0.576818	-0.790990	...	0.352996	0.8723
<b>204906</b>	135497.0	2.095278	-0.693180	-1.575512	-0.753739	-0.248853	-0.635268	-0.491066	-0.243448	-0.351194	...	0.123200	0.3245
<b>178821</b>	123796.0	-0.089165	0.807445	-0.169992	-0.959017	1.048850	-0.200202	0.876518	-0.042459	-0.201535	...	-0.263317	-0.6290
<b>68959</b>	53206.0	1.316849	-0.505091	0.660849	-0.788016	-0.970255	-0.376154	-0.728547	0.076743	-1.054445	...	-0.005768	-0.1273
...	...	...	...	...	...	...	...	...	...	...	...	...	...
<b>279863</b>	169142.0	-1.927883	1.125653	-4.518331	1.749293	-1.566487	-2.010494	-0.882850	0.697211	-2.064945	...	0.778584	-0.3191
<b>280143</b>	169347.0	1.378559	1.289381	-5.004247	1.411850	0.442581	-1.326536	-1.413170	0.248525	-1.127396	...	0.370612	0.0282
<b>280149</b>	169351.0	-0.676143	1.126366	-2.213700	0.468308	-1.120541	-0.003346	-2.234739	1.210158	-0.652250	...	0.751826	0.8341
<b>281144</b>	169966.0	-3.113832	0.585864	-5.399730	1.817092	-0.840618	-2.943548	-2.208002	1.058733	-1.632333	...	0.583276	-0.2692
<b>281674</b>	170348.0	1.991976	0.158476	-2.583441	0.408670	1.151147	-0.096695	0.223050	-0.068384	0.577829	...	-0.164350	-0.2951

1292 rows × 31 columns



```
In [24]: new_dataset['Class'].value_counts()
```

```
Out[24]: Class
0      800
1      492
Name: count, dtype: int64
```

## Train the model




```
In [25]: X = new_dataset.drop(columns='Class', axis=1)
Y = new_dataset['Class']
```

```
In [27]: X
```

```
Out[27]:
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V20	V
<b>256222</b>	157591.0	-2.456527	2.457377	-0.278652	-1.334673	-0.388679	-0.391650	-0.056992	0.619518	1.626584	...	0.712993	-0.3803
<b>221908</b>	142745.0	2.192551	-0.697003	-2.172062	-0.869045	-0.185080	-1.774343	0.265166	-0.576818	-0.790990	...	-0.031527	0.3529
<b>204906</b>	135497.0	2.095278	-0.693180	-1.575512	-0.753739	-0.248853	-0.635268	-0.491066	-0.243448	-0.351194	...	0.234644	0.1232
<b>178821</b>	123796.0	-0.089165	0.807445	-0.169992	-0.959017	1.048850	-0.200202	0.876518	-0.042459	-0.201535	...	0.000569	-0.2633
<b>68959</b>	53206.0	1.316849	-0.505091	0.660849	-0.788016	-0.970255	-0.376154	-0.728547	0.076743	-1.054445	...	0.030210	-0.0057
...	...	...	...	...	...	...	...	...	...	...	...	...	...
<b>279863</b>	169142.0	-1.927883	1.125653	-4.518331	1.749293	-1.566487	-2.010494	-0.882850	0.697211	-2.064945	...	1.252967	0.7785
<b>280143</b>	169347.0	1.378559	1.289381	-5.004247	1.411850	0.442581	-1.326536	-1.413170	0.248525	-1.127396	...	0.226138	0.3706
<b>280149</b>	169351.0	-0.676143	1.126366	-2.213700	0.468308	-1.120541	-0.003346	-2.234739	1.210158	-0.652250	...	0.247968	0.7518
<b>281144</b>	169966.0	-3.113832	0.585864	-5.399730	1.817092	-0.840618	-2.943548	-2.208002	1.058733	-1.632333	...	0.306271	0.5832
<b>281674</b>	170348.0	1.991976	0.158476	-2.583441	0.408670	1.151147	-0.096695	0.223050	-0.068384	0.577829	...	-0.017652	-0.1643

1292 rows × 30 columns



In [28]: Y

```
Out[28]: 256222    0
          221908    0
          204906    0
          178821    0
          68959     0
          ..
          279863    1
          280143    1
          280149    1
          281144    1
          281674    1
          Name: Class, Length: 1292, dtype: int64
```

In [29]: X\_train, X\_test, Y\_train, Y\_test = train\_test\_split(X, Y, test\_size=0.2, stratify=Y, random\_state=80)

In [30]: print(X.shape, X\_train.shape, X\_test.shape)

```
(1292, 30) (1033, 30) (259, 30)
```

## Modelbuilding

In [31]: model=LogisticRegression()

In [32]: model.fit(X\_train,Y\_train)

```
Out[32]: LogisticRegression
          LogisticRegression()
```



```
In [33]: model.coef_
```

```
Out[33]: array([[ -3.13211986e-05,  1.60897874e-01,  1.75202609e-01,  
                -4.87694883e-01,  5.18734888e-01,  1.44712143e-01,  
                -1.54600926e-01, -6.63488710e-02,  2.40804930e-02,  
                -1.54601006e-01, -3.08508230e-01,  2.31229140e-01,  
                -4.64667805e-01, -6.49880243e-02, -7.97302869e-01,  
                -1.03700161e-01, -2.05438190e-01, -1.99380275e-01,  
                 1.56797814e-02, -1.88022527e-02,  1.02826591e-01,  
                 1.74783658e-01,  3.55219889e-02,  1.40924217e-02,  
                 1.54015796e-04, -3.31503607e-02, -4.30602551e-02,  
                -1.71895557e-02,  6.36357850e-02,  1.54872525e-03]])
```

```
In [34]: # accuracy on training data  
X_train_prediction = model.predict(X_train)  
training_data_accuracy = accuracy_score(X_train_prediction, Y_train)
```

```
In [35]: print('Accuracy on Training data : ', training_data_accuracy*100)
```

Accuracy on Training data : 94.09486931268151

```
In [36]: print('Accuracy on Training data : ', training_data_accuracy)
```

Accuracy on Training data : 0.9409486931268151

```
In [37]: # accuracy on testing data  
X_test_prediction = model.predict(X_test)  
testing_data_accuracy = accuracy_score(X_test_prediction, Y_test)
```

```
In [39]: print('Accuracy on Testing data : ', training_data_accuracy*100)
```

Accuracy on Testing data : 94.09486931268151

```
In [40]: print('Accuracy on Testing data : ', training_data_accuracy)
```

Accuracy on Testing data : 0.9409486931268151



```
In [41]: Trainpred = model.predict(X_train)
Testpred = model.predict(X_test)
```


```
In [43]: print(sm.classification_report(Y_test, Testpred))
```

	precision	recall	f1-score	support
0	0.94	0.95	0.95	160
1	0.92	0.91	0.91	99
accuracy			0.93	259
macro avg	0.93	0.93	0.93	259
weighted avg	0.93	0.93	0.93	259

## AUC\_ROC Curve



```
In [44]: roc = model.predict_proba(X_train)[:,-1]
fpr,tpr,threshold = sm.roc_curve(Y_train,roc)
plt.plot([0,1],[0,1], 'k--')
plt.plot(fpr, tpr, label='logistic')
plt.ylabel('TPR')
plt.xlabel('FPR')
plt.show()
```




```
In [46]: sm.roc_auc_score(Y_train, roc)
```

```
Out[46]: 0.9673942430025445
```





```
In [45]: roc1 = model.predict_proba(X_test)[: ,1]
fpr,tpr,threshold = sm.roc_curve(Y_test,roc1)
plt.plot([0,1],[0,1], 'k--')
plt.plot(fpr, tpr, label='logistic')
plt.ylabel('TPR')
plt.xlabel('FPR')
plt.show()
```



```
In [48]: sm.roc_auc_score(Y_test, roc1)
```

```
Out[48]: 0.9758838383838384
```



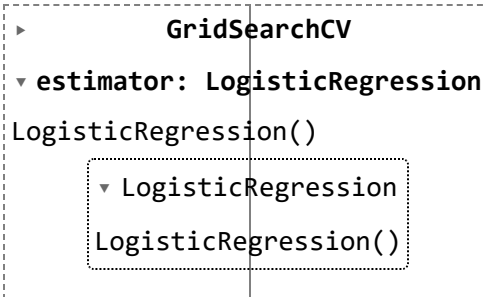
## Grid search

```
In [49]: param_grid = {  
    'penalty': ['l1', 'l2'],  
    'C': [0.1, 0.5, 1, 5, 10],  
    'solver': ['liblinear']  
}
```

```
In [52]: grid = GridSearchCV(estimator=model , param_grid=param_grid, cv=5)
```

```
In [53]: grid.fit(X_train,Y_train)
```

```
Out[53]:
```



```
  ▸      GridSearchCV  
    ▾ estimator: LogisticRegression  
      LogisticRegression()  
        ▾ LogisticRegression  
          LogisticRegression()
```

```
In [54]: best_param = grid.best_params_  
best_model = grid.best_estimator_
```

```
In [55]: y_pred = best_model.predict(X_test)
```

```
In [56]: acc = accuracy_score(Y_test,y_pred)  
pre = precision_score(Y_test,y_pred)  
rec = recall_score(Y_test,y_pred)  
f1 = f1_score(Y_test,y_pred)  
roc_auc = roc_auc_score(Y_test,y_pred)
```



```
In [57]: print(f'''  
Best param :{best_param}  
Accuracy: {acc}  
Precision: {pre}  
Recall: {rec}  
f1 score: {f1}  
AUC-ROC: {roc_auc}  
''')
```

```
Best param :{'C': 0.5, 'penalty': 'l1', 'solver': 'liblinear'}  
Accuracy: 0.9536679536679536  
Precision: 0.9578947368421052  
Recall: 0.9191919191919192  
f1 score: 0.9381443298969072  
AUC-ROC: 0.9470959595959595
```

In [ ]:

In [ ]:

